

## Concept questions

1. Why is it important to subtract the row-wise maximum before softmax?
2. In multiclass cross-entropy, why should `dlogits` be divided by batch size rather than by number of classes?
3. Why is the ReLU backward mask based on whether the pre-activation was positive?
4. Why is comparing manual gradients to PyTorch autograd a strong debugging technique?
5. Why might the highest-loss examples be good candidates for label noise review?